

Report

June 29, 2019

1 Learning Algorithm

The learning algorithm used to train this agent was DQN(Deep-Q Network).

It follows this sequence:

Learning

- Obtain random minibatch of tuples (s, a, r, s') from D
- Set target $y_i = r + \gamma \max_{a'} q(s', a, w_-)$
- Update weights $(y_i - q(s, a, w)) * \text{gradient}(w) q(s, a, w)$
- Every c steps $w \leftarrow w_-$

2 DQN Agent:

- Action Value model:

```
Input: state_size (int): 37
Output: action_size (int): 4
Layer 1 - fc1_units (int), Number of nodes in first hidden layer: 64
Activation Layer 1: RELU
Layer 2 - fc2_units (int): Number of nodes in second hidden layer: 64
Activation Layer 2: RELU
```

- Agents Hyper Parameter:

```
BUFFER_SIZE = int(1e5) # replay buffer size
BATCH_SIZE = 64        # minibatch size
GAMMA = 0.99           # discount factor
TAU = 1e-3             # for soft update of target parameters
LR = 5e-4              # learning rate
UPDATE_EVERY = 4       # how often to update the network
```

3 Plot of Rewards

Episode 100 Average Score: 0.52

Episode 200 Average Score: 3.25

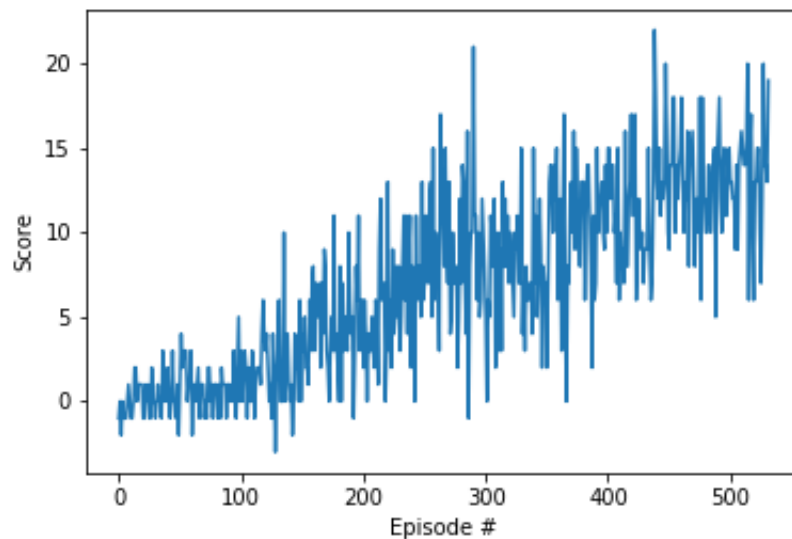
Episode 300 Average Score: 7.45

Episode 400 Average Score: 9.02

Episode 500 Average Score: 12.30

Episode 532 Average Score: 13.04

Environment solved in 432 episodes! Average Score: 13.04



BaseScores.png

4 Ideas for Future Work

It would be interesting to try the following modifications:

- Double DQN (DDQN) Prioritized experience replay Dueling DQN

- It would also be interesting to try learning directly from the pixels